

Project Report: Conflict-Free Real-Time Workspace

Project Name: The Conflict-Free Diagramming Engine

Architecture: MERN Stack (MongoDB, Express, React, Node.js) + WebSockets

Developer Tier: Expert / Industry-Ready

1. Executive Summary

This project is a high-performance, real-time collaborative workspace designed to handle complex graph data (Mind Maps/Diagrams). Unlike standard CRUD applications, this system implements **Distributed Systems logic** to handle offline synchronization, high-frequency updates, and deep hierarchical data structures without performance degradation.

2. Technical Challenge Breakdown

Challenge A: High-Performance Graph Storage

The Problem: Storing 1,000+ deeply nested nodes in a document-based database (MongoDB).

The Solution: * **Materialized Path Pattern:** Each node stores its lineage in a string field (**path**). This allows for a single Regex-based query to fetch an entire subtree, bypassing the need for expensive recursive lookups.

- **Indexing:** Heavily indexed paths ensure $\$O(n)\$$ or better retrieval speeds for large datasets.

Challenge B: Offline Conflict Resolution

The Problem: Data loss occurring when a user performs actions on a deleted parent while offline.

The Solution: * **Lost & Found Logic:** Instead of a "Hard Delete" or "Silent Failure," the server validates parent existence upon reconnection.

- **Orphan Management:** If a parent is missing, the system automatically reparents the orphan nodes to a specific "Lost & Found" root, preserving user work while maintaining data integrity.

Challenge C: Frontend Virtualization (10,000+ Items)

The Problem: Browser DOM exhaustion when rendering thousands of interactive SVG/HTML elements.

The Solution: * **Viewport Culling:** A custom React hook calculates the intersection of node coordinates and the current viewport (considering zoom and pan).

- **GPU Acceleration:** Using `transform: translate3d()` to offload canvas movement to the GPU, keeping the interface at a steady 60fps.

Challenge D: Temporal Data Replay (Time Travel)

The Problem: Traditional "Undo" is destructive in a collaborative environment.

The Solution: * **Event Sourcing:** Every mutation (move, rename, delete) is stored as a discrete event in a `ChangeLog` collection.

- **State Reconstruction:** To view the workspace at "Yesterday, 4 PM," the system fetches a snapshot and replays all log entries up to that specific timestamp in memory.

3. System Architecture

Layer	Technology	Key Responsibility
Frontend	React (Hooks + Context)	Optimistic UI updates, Viewport Math, WebSocket Client.
Transport	WebSockets (<code>ws</code> library)	Low-latency, bi-directional binary/JSON data transfer.
Backend	Node.js / Express	Conflict validation, Broadcast logic, ChangeLog generation.
Database	MongoDB	Materialized Path storage, Operational logging.

4. Key Performance Metrics

- **Latency:** Under 50ms for local broadcast (Optimistic UI makes it feel like 0ms).
- **Scalability:** Tested for 1,000+ nodes via the `seed.js` script.

- **Efficiency:** WebSocket throttling limits database writes to 20 times per second per user, preventing server flooding.
-

5. Conclusion & Future Roadmap

The current implementation successfully solves the "Offline Nightmare" and the "10k Node Freeze." The architecture is modular and ready for production deployment.

Future Enhancements:

1. **CRDT Integration:** Upgrading from LWW (Last Write Wins) to Yjs for character-by-character text editing inside nodes.
2. **Binary Serialization:** Transitioning from JSON to Protocol Buffers to further reduce bandwidth by 60%.
3. **Presence Indicators:** Implementing "Live Cursors" to show where other collaborators are looking.